

Securing Internet communications



License & Disclaimer

2

License Information

This presentation is licensed under the
Creative Commons BY-NC License



To view a copy of the license, visit:

<http://creativecommons.org/licenses/by-nc/3.0/legalcode>

Disclaimer

- We disclaim any warranties or representations as to the accuracy or completeness of this material.
- Materials are provided “as is” without warranty of any kind, either express or implied, including without limitation, warranties of merchantability, fitness for a particular purpose, and non-infringement.
- Under no circumstances shall we be liable for any loss, damage, liability or expense incurred or suffered which is claimed to have resulted from use of this material.

Topics

3

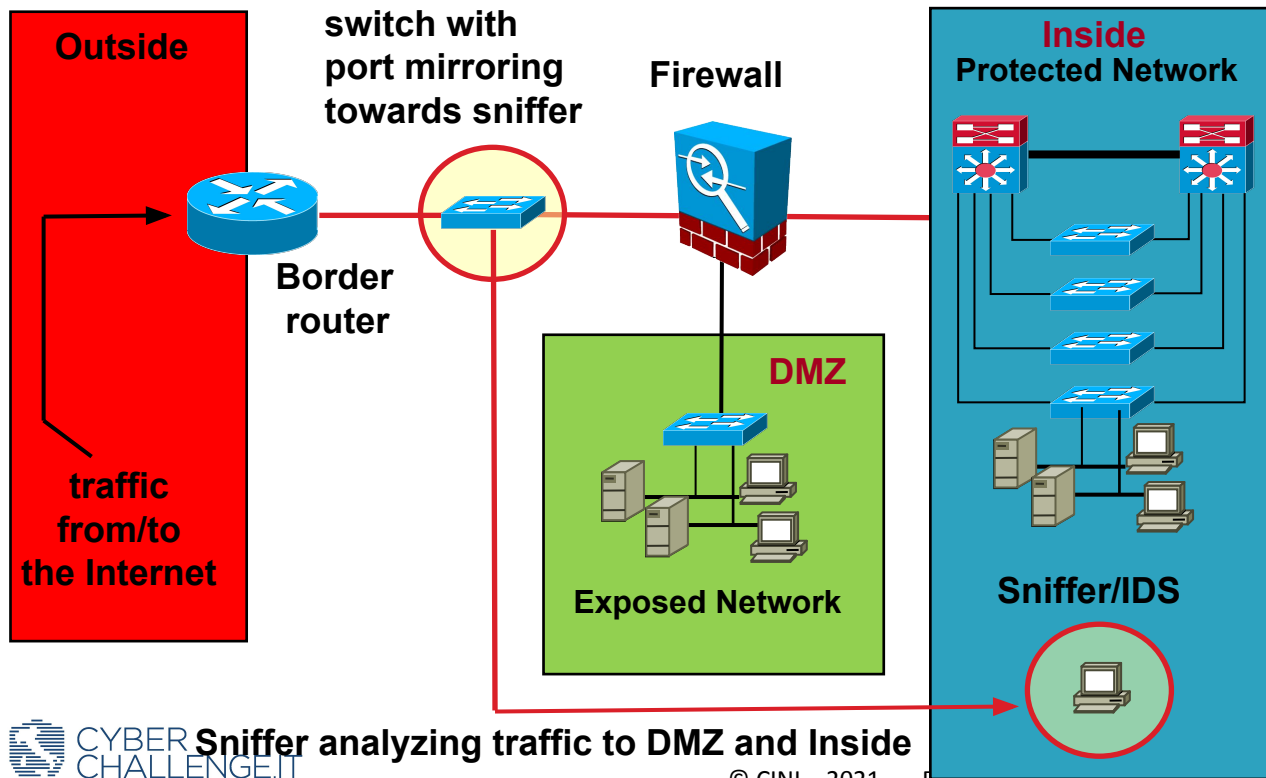
- Security enforcement devices
- Traffic filtering strategies and policies
- Network access control: techniques and tools
- Implementing simple network access control policies

Current Topic

4

- Security enforcement devices
- Traffic filtering strategies and policies
- Network access control: techniques and tools
- Implementing simple network access control policies

Basic security architecture



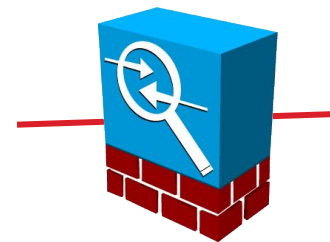
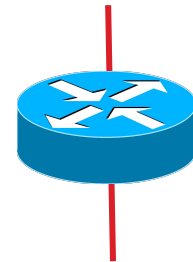
□ In a common network architecture there are at least three domains:

- **Outside** (all the world outside - the Internet): trust degree 0
- **Inside** (the internal organization to be protected and hidden): degree of trust 100
- **DMZ** (the set of internal machines that expose services outside): degree of trust $0 < x < 100$

Security enforcement devices

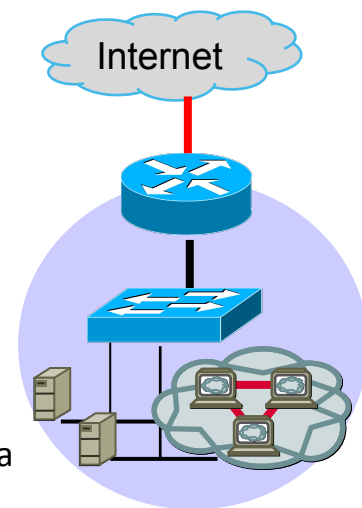
6

- In a common security architecture in order to implement security policies we can rely on two perimeter control devices:
 - **Routers** (typically those located on the network border)
 - **Firewalls**



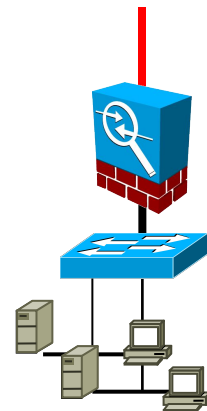
Border Router

- ❑ The border router is the first barrier protecting its internal network
 - ❑ difficult to circumvent by malicious end-users
- ❑ It allows the centralization of a good number of security checks
- ❑ Its protection is fundamental
 - ❑ a compromise may open access to the internal network
 - ❑ an inadequate filtering policy can expose the internal network to attacks
 - ❑ corruption of routing tables can cause disruptions and unauthorized access to data
- ❑ A properly configured router can minimize effects from internal sites compromised by attacks or hostile activities (insider threats)



Firewall

- ❑ Firewall is an english term with original meaning of a fire isolation barrier
- ❑ It is the main passive perimeter defense component
- ❑ It has security enforcement tasks, in the broadest sense of the term, with the aim of controlling traffic between two or more networks:
 - ❑ allowing only what is specifically authorized by the security policy
 - ❑ detecting and reporting any attempts to violate the security policy
 - ❑ possibly carrying out additional auditing and accounting functions
 - ❑ it can also connect at the link or network layer two or more network segments



Why installing a firewall

- To implement a security policy:
 - Able to allow controlled access to systems or services of a protected network:
 - Authorized users only
 - Only to authorized systems
 - Able to allow users and systems of a protected network to access the systems and services of an outside (untrusted) network in a controlled way:
 - only if the risk is acceptable
 - recording all their activities



Firewall: pros

- ❑ Centralization of security policies
 - ❑ Can result in a Single point of failure (can be a disadvantage)
- ❑ Relying on a special purpose solution able to optimize traffic filtering operations (through appropriate HW)
- ❑ Ability to inspect traffic from data link to application layers
- ❑ Stateful control of sessions



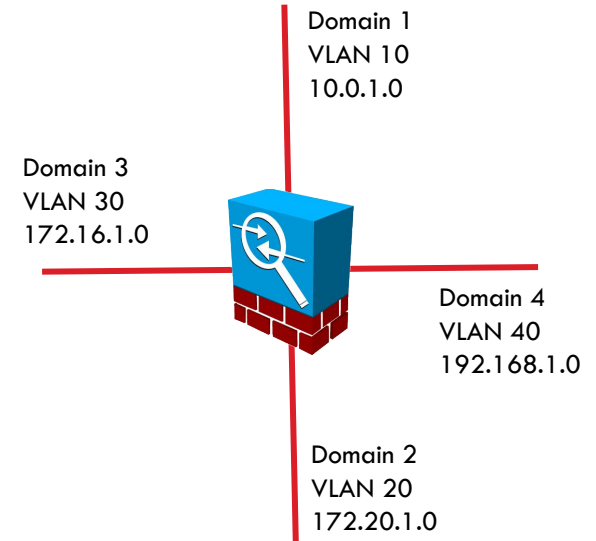
Firewall: cons

- ❑ Difficulty in coping with non-trivial protocols
- ❑ Performance / throughput
 - ❑ It can turn into a bottleneck
 - ❑ user perception can be negative due to service limitation
- ❑ Complex management
 - ❑ Configuration requires specialization
 - ❑ Verification and analysis of logs is not straightforward
- ❑ Excessive sense of trust and internal insecurity
- ❑ High costs for performance beyond Gigabit



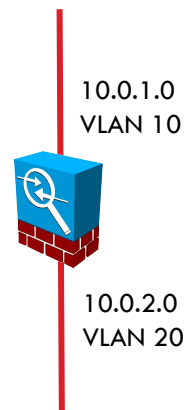
Implementation and basic functions

- ❑ Network device with at least 2 network interfaces
 - ❑ Each interface identifies a **separate security domain** on a different network segment (VLAN)
- ❑ Can remap IP addresses (NAT)
- ❑ Filters traffic between different zones/domains through predefined rules (access control policies)
- ❑ It can mediate access to specific applications for control and inspection purposes:
 - ❑ Proxy service access
 - ❑ Content filtering (selective content filtering)
 - ❑ Deep packet inspection and traffic analysis
 - ❑ Enforce bandwidth limitations on specific traffic types

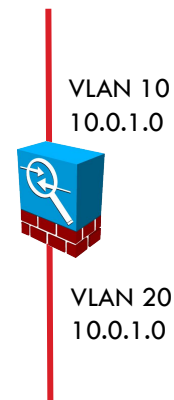


Firewall: operating modes

- A firewall can operate in two ways:
 - **Routed**: Operates at level 3, segmenting different networks based on IP addresses
 - **Transparent**: Operates at level 2, segmenting on MAC address basis
- A routed firewall looks like a layer 3 device and needs an IP address/network on each interface associated to a segment
 - Routes IP/IPv6 traffic between the various interfaces
 - Supports the most common routing protocols



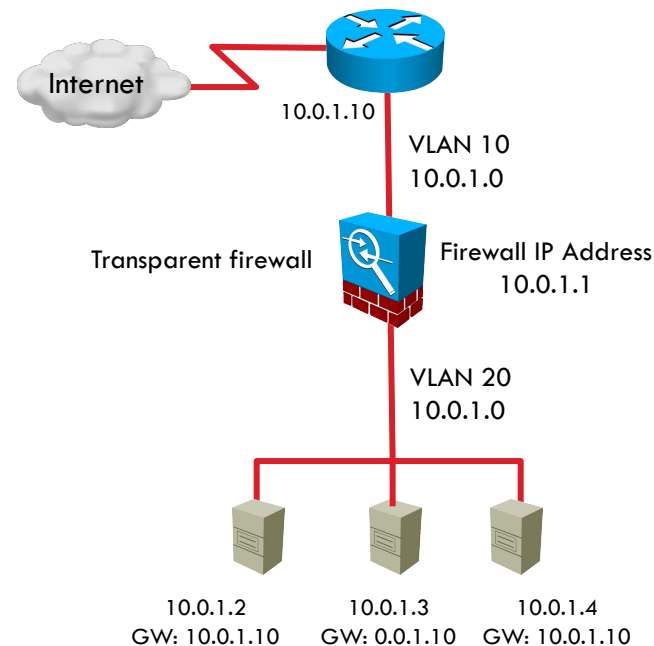
Routed mode
or L3 firewall



Transparent mode
or L2 firewall

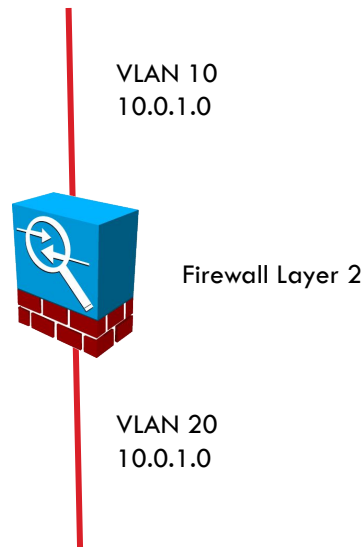
Transparent mode operations

- Layer 3 traffic must be explicitly allowed to pass through the firewall
 - However it performs packet screening/filtering from network to application layers
- The segments connected to the interfaces must be on the same layer 3 subnet
- The firewall IP address must not be configured as the default gateway for connected devices
 - Devices must point to the router ahead of the firewall (passed through transparently)
 - Each interface identifies a different segment/VLAN even if associated with the same IP network (the firewall bridges different segments)



Transparent mode benefits

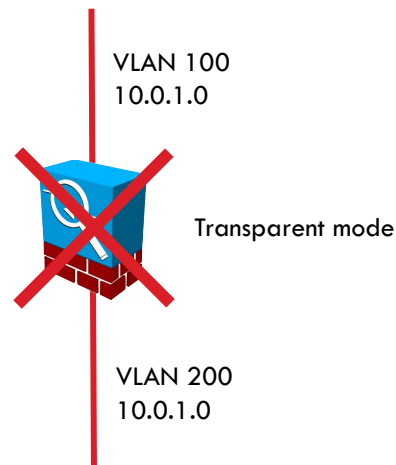
- Flexible, integrated and easy to manage:
 - IP-level redirection not required
 - No NAT to configure
 - Routing and redirection problems cannot occur (it does not perform routing)
- Totally invisible from the outside
- Greater robustness



Transparent mode unsupported features

□ The following features are typically not supported by a firewall in transparent mode:

- NAT
- Routing protocols (e.g. OSPF, RIP, BGP)
- IP / IPv6
- DHCP relay
- QoS
- Multicast
- VPN termination



Current Topic

17

- Security enforcement devices
- **Traffic filtering strategies and policies**
- Network access control: techniques and tools
- Implementing simple network access control policies

Traffic filtering: putting controls on border router

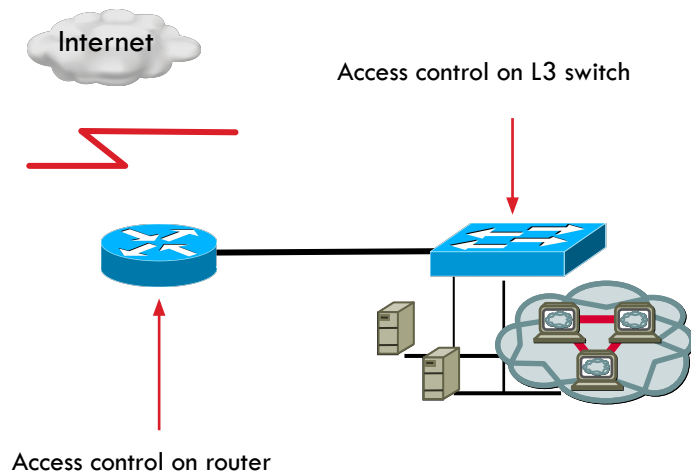
- ❑ Router and Switch Layer 3 devices provide simple access control mechanisms

- ❑ Based on stateless traffic filtering
 - ❑ Filters based on IP address and TCP/UDP Ports

- ❑ The use of complex controls with a significant number of filtering clauses, entails a certain increase in CPU load in the forwarding activity.

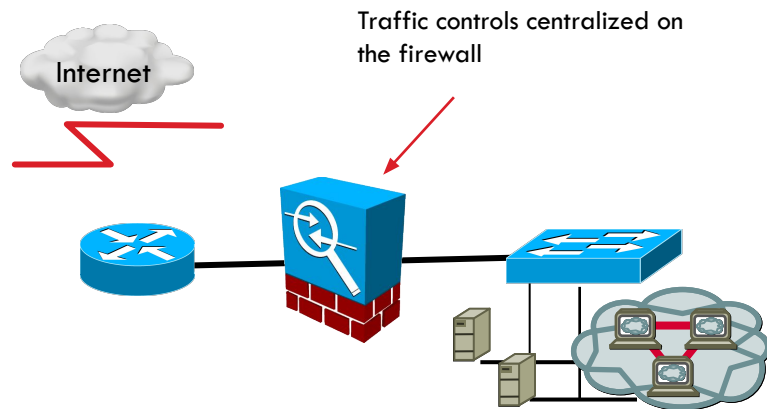
- ❑ can be acceptable if the Router or L3 Switch provide hardware implementation of access control mechanisms

- ❑ The truly advantage is that such devices are already present in any network, partitioning it in a natural way



Traffic filtering: putting controls on firewalls

- ❑ The introduction of a firewall reduces the CPU burden associated to traffic control/filtering activity on routers or L3 switches
- ❑ The centralization of control policies on the firewall constitutes a significant advantage from the management perspective:
 - ❑ reduces the configuration complexity
 - ❑ centralizes the management of logics and filtering problems
 - ❑ It allows you to simultaneously protect thousands of machines
- ❑ This policy does not scale in the presence of large traffic volumes and becomes a performance bottleneck that can be exploited to create DoS



Choosing the right filtering location

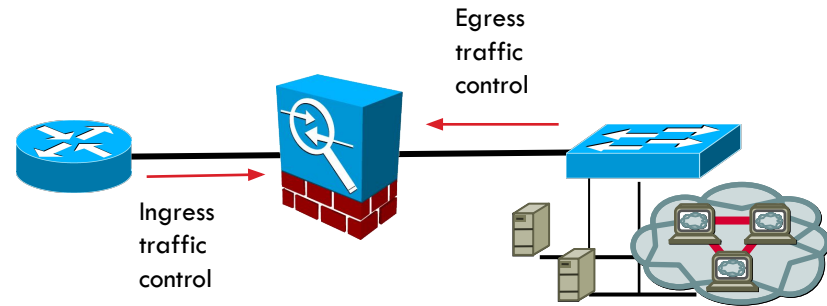
❑ Ingress filtering:

- ❑ natively knows from which interface packets are coming in
- ❑ provides protection of internal networks

❑ Egress filtering:

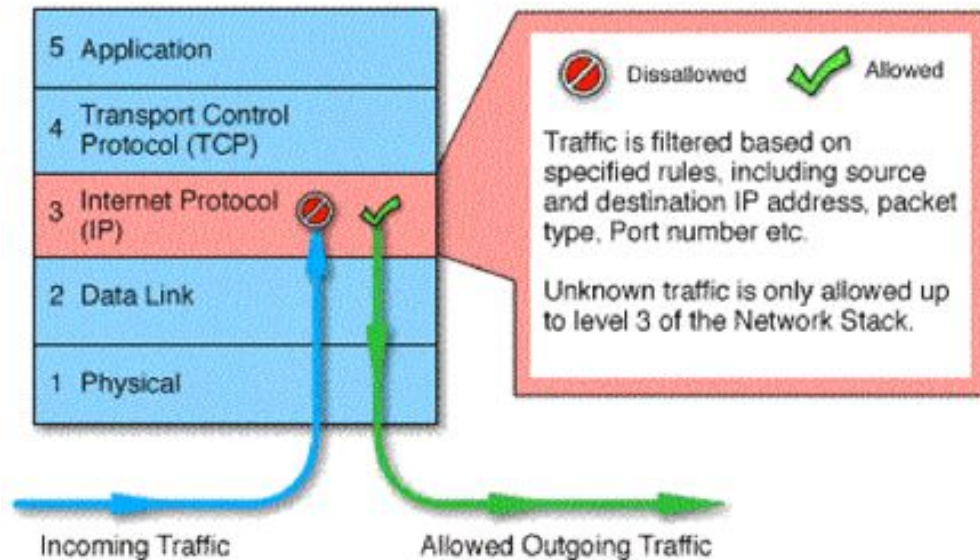
- ❑ enables checking also the locally generated traffic exiting the domain
- ❑ blocks what should not come out

❑ The controls should be located **as close as possible** to the traffic origin



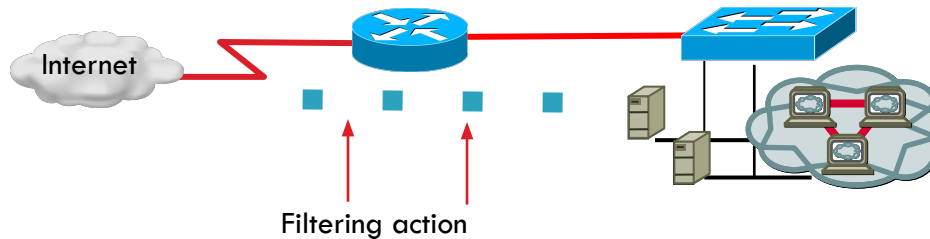
Filtering Parameters

- IP Header
 - source address
 - destination address
 - protocol
 - flags, options (source routing)
- TCP/UDP Header
 - source port
 - destination ports
 - flags TCP (SYN, ACK)



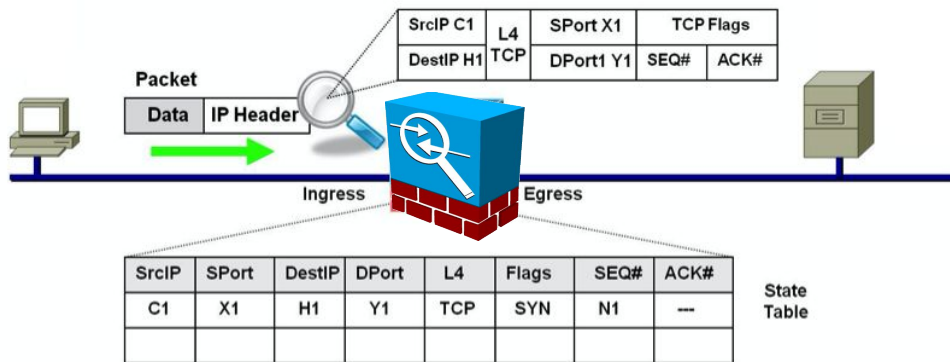
Stateless filtering (router)

- ❑ Based only on IP addresses, TCP / UDP ports (source and destination), and protocol
- ❑ Controls are carried out independently one packet at a time (no memory/state)
- ❑ There is no perception of the flow of packets belonging to an end to end connection
- ❑ Packets can also come from interfaces other than the one on which they exit (asymmetric forwarding phenomena are tolerated)



Stateful filtering (firewall)

- When a new connection is established, if the filtering rules do not block it, then the related information are used to add an entry (**session**) of a **connection status table**.
- Subsequent incoming packets will be handled according to their belonging to one of the active connections (or data flow sessions) whose status is saved in the table.
- When the connection is terminated, the corresponding entry in the table is deleted
- The table contains:
 - Unique session ID
 - Connection status (*handshaking, established, closing*)
 - Packet sequencing information
 - Source and destination addresses/ports
 - Network interfaces used



Filtering (access control) policies

- Before defining any filtering policy aimed at performing access control, a careful preliminary assessment must be made, by considering:
 - Who needs access?
 - When and how?
 - From where?
 - At which time/date?
 - What services does it need?
 - What protocols does it use?
 - What QoS (e.g bandwidth) does it require?



Filtering (access control) policies

A firewall (or router) can operate in two diametrically opposite ways:

❑ **Deny All:** Anything that isn't specifically allowed is denied. **High security**

- ❑ Block all traffic and each service must be implemented on a case-by-case basis
- ❑ More conservative policy in terms of protection
- ❑ the number of choices available to the user is limited

❑ **Allow All:** Anything that is not specifically denied is allowed. **Ease of management**

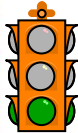
- ❑ Forward all traffic and each malicious service must be closed on a case-by-case basis
- ❑ Increasing difficulties in guaranteeing security as the network grows.
- ❑ Rarely used in security schemes, however it is may cover several specific cases



Selective traffic filtering

26

Service	Port	Protocol
echo	7	TCP/UDP
discard	9	TCP/UDP
systat	11	TCP/UDP
daytime	13	TCP/UDP
netstat	15	TCP
quotd	17	TCP/UDP
chargen	19	TCP/UDP
ftp-data	20	TCP
ftp	21	TCP
ssh	22	TCP/UDP
telnet	23	TCP
smtp	25	TCP
time	37	TCP/UDP
rip	39	TCP/UDP
whois	43	TCP/UDP
tacacs	49	TCP/UDP
domain	53	TCP
whois++	63	TCP/UDP
bootp	67-68	UDP
tftp	69	UDP
gopher	70	TCP
finger	79	TCP
http	80	TCP
link	87	TCP
supdup	95	TCP
pop2	109	TCP
pop3	110	TCP
sunrpc	111	TCP/UDP
auth	113	TCP/UDP
nnntp	119	TCP
ntp	123	TCP/UDP
nbios-ns	137	TCP/UDP

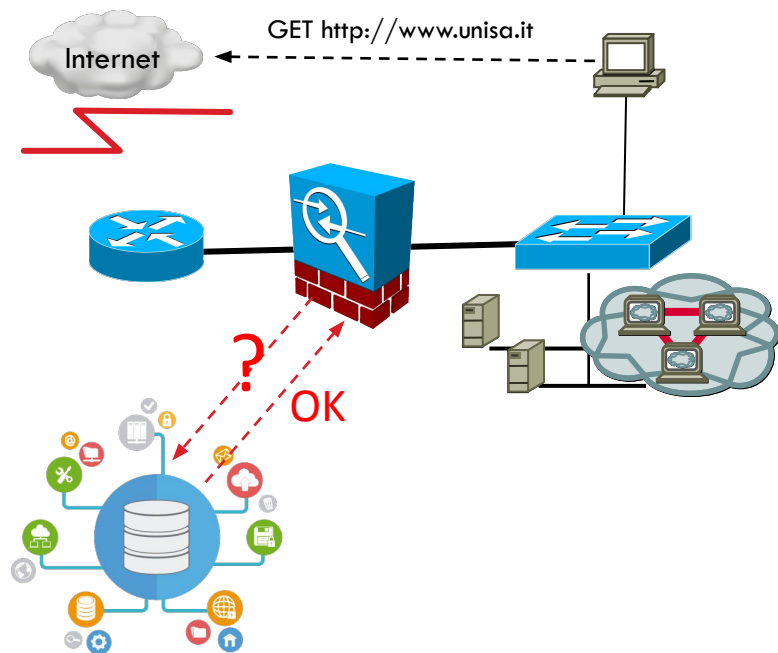


Service	Port	Protocol
nbios-dgm	138	TCP/UDP
nbios-ssn	139	TCP/UDP
imap	143	TCP
NeWS	144	TCP
snmp	161	UDP
snmptrap	162	UDP
xmcp	177	UDP
irc	194	TCP/UDP
weis/Z39.50	210	TCP
imap3	220	TCP
ldap	389	TCP/UDP
network-ip	396	TCP/UDP
rmt	411	TCP
https	443	TCP
exec	512	TCP
biff	512	UDP
login	513	TCP
who	513	UDP
shell	514	TCP
syslog	514	UDP
printer	515	TCP/UDP
talk/intalk	517-518	TCP/UDP
route	520	UDP
timed	525	TCP/UDP
uucp	540-541	TCP
mountd	635	TCP/UDP
wins	1512	TCP/UDP
radius-old	1645-1646	UDP
radius	1812-1813	UDP
openwin	2000	TCP
NFS	2049	TCP/UDP
X11	6000-6063	TCP

- It can be a good practice blocking or selectively filtering potentially dangerous services
- Only allow access to an extremely limited number of services (e-mail, www, ftp) provided by specific and possibly controlled hosts

Content filtering

27



- ❑ Filtering unwanted, objectionable, and harmful content through URL inspection
- ❑ Requires the use of third party and always up-to-date knowledge bases
 - ❑ Resource classification DB
 - ❑ Categorization engines
- ❑ The firewall performs payload inspection and before admitting the session checks the content type against local policies
 - ❑ e.g. block gambling, drugs, crime-related URLs

Current Topic

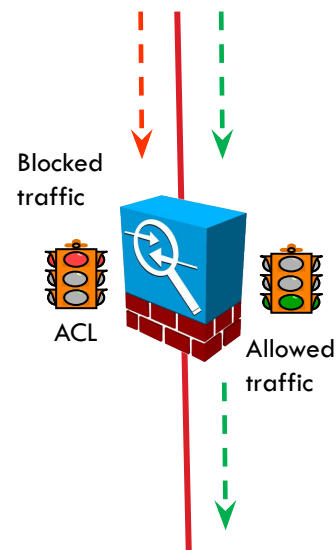
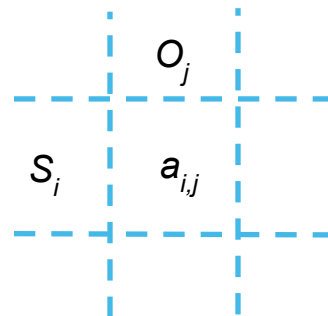
28

- Security enforcement devices
- Traffic filtering strategies and policies
- **Network access control: techniques and tools**
- Implementing simple network access control policies

Traffic filtering mechanisms: ACL

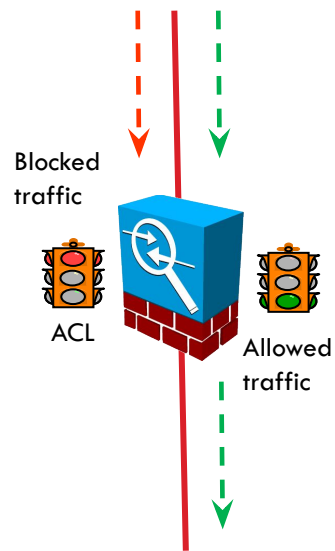
- The simplest and most immediate way to implement security schemes and policies is traffic filtering (packet filtering)
- Routers, switches and firewalls support lists of filtering (or access control) rules known as ACLs: **Access control Lists**
- An ACL represents a column in the Lampson's **access control matrix**, where:

- S_i : i -th subject to be controlled (e.g. a netblock)
- O_j : j -th object to be protected (e.g. an interface)
- a_{ij} : access rights of S_i on O_j (e.g. permit or deny)



Traffic filtering mechanisms: ACL

- Filtering rules can be applied at:
 - **data link** layer (based on MAC addresses)
 - **network** layer (based on IP addresses)
 - **transport** layer (based on ports or protocol)
- Additional elements may be checked:
 - Date and time of application
 - Session flags or status (established, closing etc.)
- Each packet received is compared with each rule, in the order in which it appears on the list, to decide if it has to be forwarded or dropped
 - The application of controls takes place on an interface basis
 - Eligible actions are **permit** (or allow) and **deny** (or drop)
 - The direction of application of the controls (**inbound** or **outbound**) is significant and defines the origin of the traffic concerned

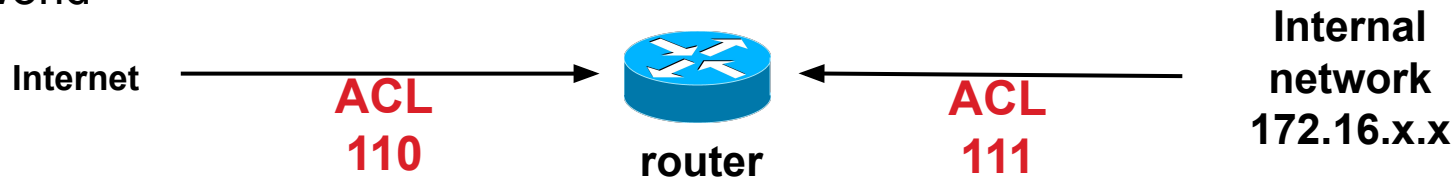


ACL Syntax

- Only one ACL can be applied to an interface in each specific direction (inbound or outbound):

```
interface ethernet 0
  ip access-group 110 in
  ip access-group 111 out
```

- In the example below, the ACL 110 and 111 are applied respectively to the input and output on the border interface that connects a router to the outside world



ACL Syntax

- An ACL is made up of rules scanned in sequence until the first match

The diagram illustrates the structure of an access-list entry. It is divided into three main sections: list identifier, protocol, and destination address/port. The list identifier section contains the text 'access-list 110'. The protocol section contains 'permit tcp'. The destination address/port section contains 'any gt 1024 172.16.0.0 0.0.255.255 eq 23'. Brackets are used to group the components: 'permit' is grouped under 'action', 'tcp' is grouped under 'protocol', 'any gt 1024' is grouped under 'src addr/port', and '172.16.0.0 0.0.255.255 eq 23' is grouped under 'destination address/port'.

list identifier	protocol	destination address/port
access-list 110	permit tcp	any gt 1024 172.16.0.0 0.0.255.255 eq 23
	action	src addr/port

- The masks associated with the addresses are in "reverse dotted mask" format or "/msklen" in format (e.g. 0.0.0.255 is equal to /24)

ACL Syntax

- ❑ **Long search**: the search is carried out until there is a matching (rule with permission or denial found) or until the list is finished;
- ❑ Efficiency depends on the **order**: the most frequent matching element should be the first in the list
- ❑ Removing a permission may be without effect
- ❑ The **any** option replaces 0.0.0.0 as the IP address and 255.255.255.255 as the wildcard mask. It results in a matching with any compared address.
- ❑ The **host** option replaces 0.0.0.0 as a mask. This mask requires that all the bits of the address match. Compare exactly one address.

ACL Syntax

- Every ACL terminates with an implicit “deny any any” clause
- It is possible to use relational operators in ACLs: eq neq, gt, lt:

```
access-list 110 deny tcp 192.168.1.0 0.0.0.255 any eq www
access-list 110 deny tcp any eq ftp 192.168.1.25
```

- ACLs can be assigned logical names

```
ip access-list extended allowt permit tcp host 192.132.34.17 any eq 23
```

- It is possible to define rules in ACL that can be activated on a date/time basis, specifying a "time-range" of validity and a periodic or absolute scope

```
time-range no-http periodic weekdays 8:00 to 18:00
access-list 110 deny tcp any any eq http time-range no-http
```

ACL Syntax

- The "established" clause at the end of a rule identifies all TCP connections that have passed the setup phase (3 way handshake)

```
access-list 110 permit tcp any any established
```

- allows you to block all incoming traffic from the outside, with the exception of return TCP traffic, due to a TCP session started from the inside.
- checks, on incoming TCP packets, the presence of the TCP ACK or RST flags:
 - if they are present, traffic is allowed,
 - otherwise it is assumed that the traffic has been generated from the outside and will be blocked.

ACLs on Linux: iptables

- Simple ACLs can be implemented as well under linux with iptables
- **iptables** is used to set up, maintain, and inspect the tables of IPv4 packet filter rules in the Linux kernel.
- Several different tables may be defined. Each table contains a number of built-in chains and may also contain user-defined chains.
- **chain** = list of **rules** which can match a set of packets
 - each rule specifies criteria for a packet and an associated **target**, namely what to do with a packet that matches the pattern
 - We are interested in the FORWARD built-in chain:
 - packets that have been routed and were not for local delivery will traverse this chain.

ACL Syntax: iptables

- It is possible a new user-defined chain by the given name

```
iptables -N acl111
```

- ... and apply on specific inbound or outbound interfaces

```
iptables -A FORWARD -i eth1 -o eth0 -j acl110
```

- Also a default policy for the chain can be specified

```
iptables -P FORWARD DROP
```

- Possible targets are

- accept = let the packet through
- drop = drop the packet on the floor

ACL Syntax: iptables

- Syntax is very intuitive and based on traditional shell command-line

Diagram illustrating the syntax of an iptables rule:

```
iptables -A myacl -p tcp -s 172.16.1.0/24 -d 192.168.1.1 --dport domain -j ACCEPT
```

The components are labeled as follows:

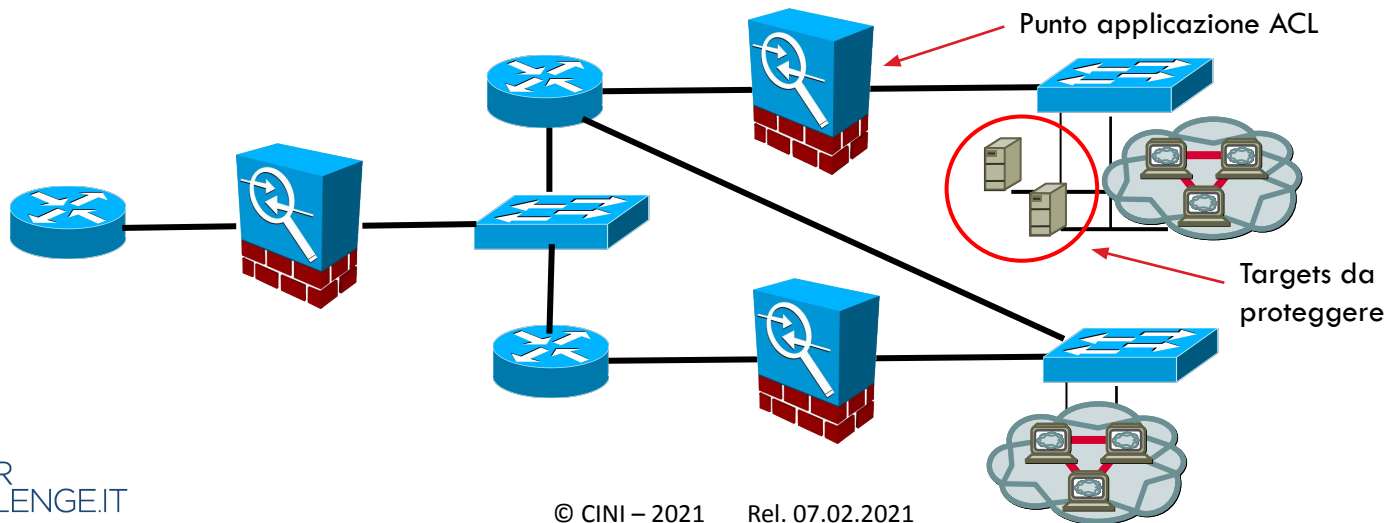
- list identifier**: myacl
- protocol**: tcp
- src addr/mask**: 172.16.1.0/24
- destination address/port**: 192.168.1.1 --dport domain
- action**: -j ACCEPT

- very similar to traditional ACLs and also provides an “established” facility

```
iptables -A acl110 -m state --state ESTABLISHED,RELATED -j ACCEPT
```

Applying ACL in the correct places

- ACLs should be placed **as close as possible** to the **target** to be protected
- This allows to **restrict** the **size** of the **security domain** in order to increase the effectiveness of the filtering policies implemented and make the solution more **scalable**



ACL applied on a switch

- Filtering can be applied also at the link layer
 - e.g. by authorizing only one host to traverse an interface

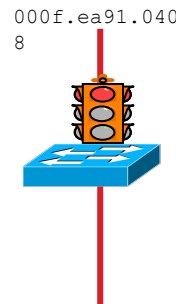
```
mac access-list mac-01          interface eth 1
    permit host 00c0.4f00.0407 any      mac port access-group mac-01 in
```

```
iptables -A FORWARD -i eth1 -m mac --mac-source 00:C0:4F:00:04:07 -j ACCEPT
```

- It can be helpful to completely lock the mac of a compromised host

```
mac-address-table static 000f.ea91.0408 vlan 1 drop
```

```
iptables -A FORWARD -m mac --mac-source 00:0F:EA:91:04:08 -j DROP
```



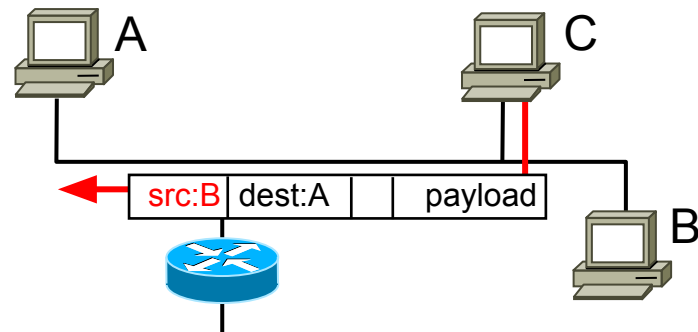
Current Topic

41

- Security enforcement devices
- Traffic filtering strategies and policies
- Network access control: techniques and tools
- **Implementing simple network access control policies**

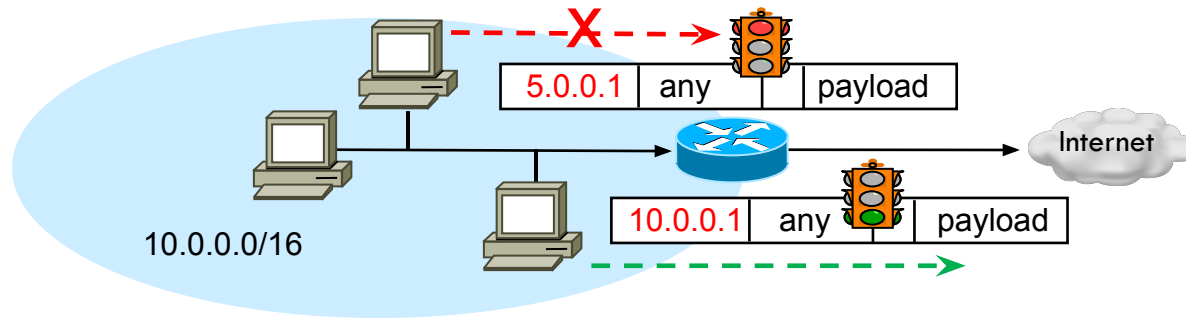
IP address spoofing

- ❑ The source IP address is currently the only mechanism for identifying the source available on the Internet
- ❑ The falsification/forgery of this data is the basis of most of the attacks and hostile actions
- ❑ Spoofing consists in falsifying the source address
 - ❑ Any user is able to generate IP packets with any value of the fields provided by the protocol structure
 - ❑ Therefore it is immediate to change the source address of the IP packets to prevent any form of identification
 - ❑ The result is that C in attacking A assumes the identity of B



Inbound anti-spoofing filtering

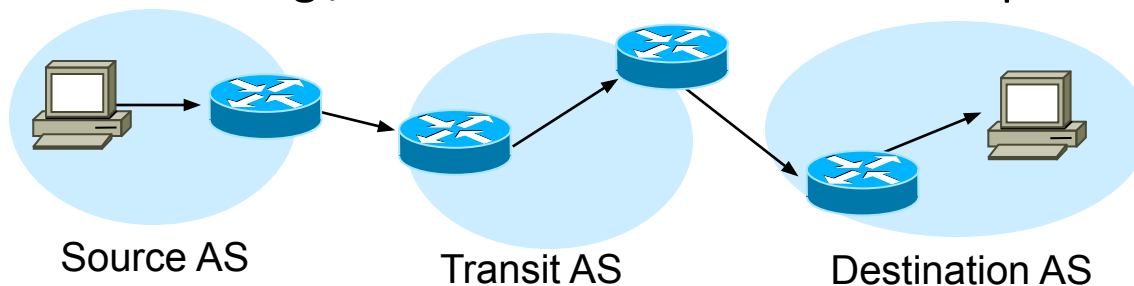
- Solution: checking and enforcing the correctness of the origin of the generated packets



- Inbound filtering policy (RFC 2827, 2000): A border router forwards only packets with legitimate source addresses

Practical implementation problems

- ❑ It is necessary that all the organizations involved and the transit ISPs do it
- ❑ Everything is based on a collaboration and trust logic working at a global level
 - ❑ If 10% of ISPs do not implement it, it is ineffective
- ❑ Another solution: enforcing / IP validation of sources at AS peering level



- ❑ A packet can only pass if the transit AS validates the source

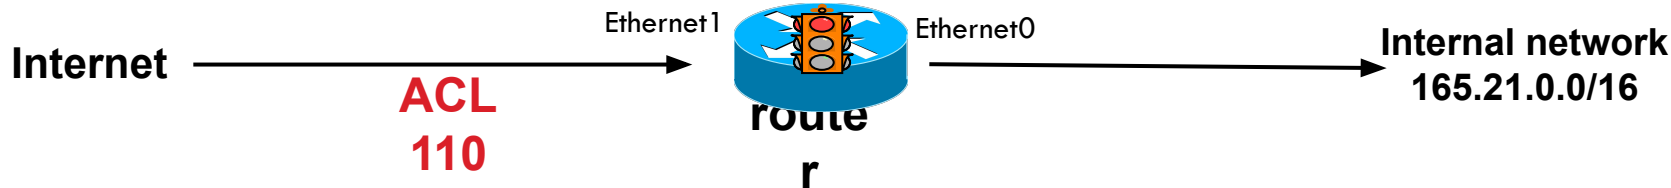
Inbound Anti-Spoofing filters

- The easiest way to protect yourself is to discard all incoming traffic with inadmissible source addresses with respect to the traffic origin

```
interface ethernet 1  
ip access-group 110 in
```

```
# blocca traffico spoof entrante da eth1  
iptables -A FORWARD -i eth1 ...
```

- Block all traffic with source addresses 165.21.0.0/16 if coming from outside (they are my internal addresses!)



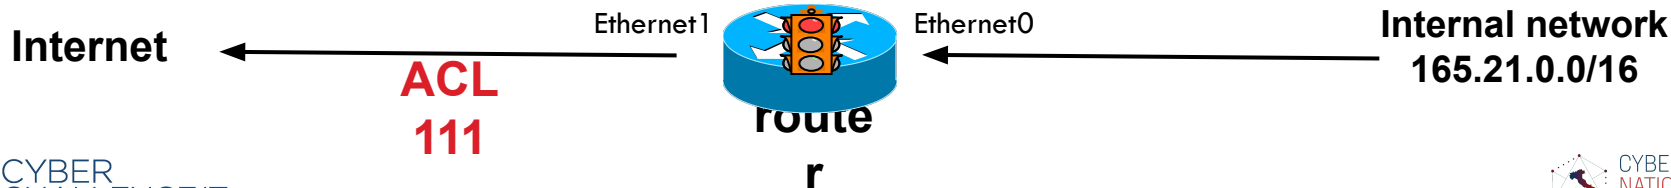
Outbound Anti-Spoofing filters

- To also prevent voluntary or involuntary spoofing from the inside of your network to the outside, similar filtering must be applied to outbound traffic

```
interface ethernet 1
ip access-group 111 out
```

```
# non inoltrare il traffico spoof da eth0
iptables -A FORWARD -i eth0 ...
```

- Block any outgoing packet with source address that does not fall on the network 165.21.0.0/16



Anti-Spoofing ACLs

Inbound Anti spoofing

! Block traffic from the outside with internal source addresses:

```
access-list 110 deny ip 165.21.0.0 0.0.255.255 any log  
access-list 110 permit ip any any
```

```
iptables -A FORWARD -i eth1 -s 165.21.0.0 /16 -j DROP
```

Outbound Anti spoofing

! Block outgoing traffic with foreign source IPs:

```
access-list 111 permit ip 165.21.0.0 0.0.255.255 any  
access-list 111 deny ip any any log
```

```
iptables -A FORWARD -i eth0 -s ! 165.21.0.0 /16 -j DROP
```

Docker Fundamentals



Outline

49

- Docker images and containers
- Docker networking
- Docker compose

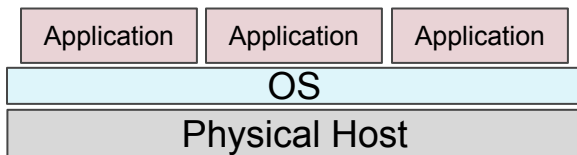
Outline

50

- Docker images and containers
- Docker networking
- Docker compose

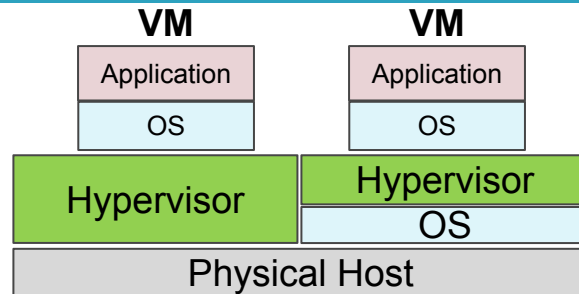
Traditional vs Virtualized Deployment

51



- Physical Hosts run an Operating System (e.g., Windows or Linux).
- Multiple applications run on the shared OS.

¹<https://www.virtualbox.org/>



- A special software, i.e., the **Hypervisor**, provides Virtual Machines.
- Examples of such technologies are *Virtualbox*¹ or *Linux KVM*².
- VM is a full machine running all the components, including its own Operating System, on top of the virtualized hardware

²<https://www.linux-kvm.org>

Traditional vs Virtualized Deployment

52

Traditional Deployment

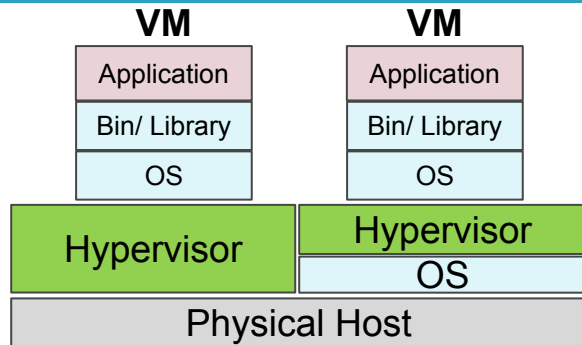
- ❑ No way to define resource boundaries for applications.
- ❑ Isolating applications requires running them on different physical servers (expensive and resources could be underutilized).

Virtualized Deployment

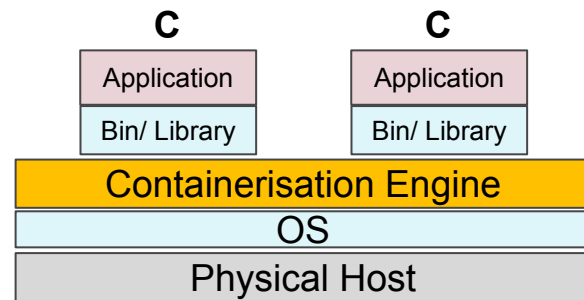
- ❑ Virtualization allows:
 - ❑ applications to be isolated between VMs
 - ❑ better utilization of resources in a physical server
 - ❑ better scalability.

Virtualized vs Container Deployment

53



- ❑ Virtual hardware
 - ❑ Each VM has an OS and Application
 - ❑ Share hardware resource from the Physical Host



- ❑ Virtual Operating Systems
 - ❑ Isolated environments, namely **containers**, sharing the same *real* operating system
 - ❑ Containers run from a distinct image that provides all files (Bin/, Library) necessary to support them
 - ❑ Examples of such technologies is *Docker*¹.

¹<https://www.docker.com/>

Virtualized vs Container Deployment

54

Virtualized Deployment

- ❑ Heavyweight: each VM relies on a full copy of an Operating System.
- ❑ Provides full isolation.
- ❑ Best suited for when you have applications that need to run on *different* Operating System flavors.

Container Deployment

- ❑ Lightweight: sharing OS resources significantly reduces the overhead required for running containers.
- ❑ Provides a (relaxed) process-level isolation.
- ❑ Best suited for when you have applications that need to run over a *single* Operating System kernel.

Docker images and containers

55

Image

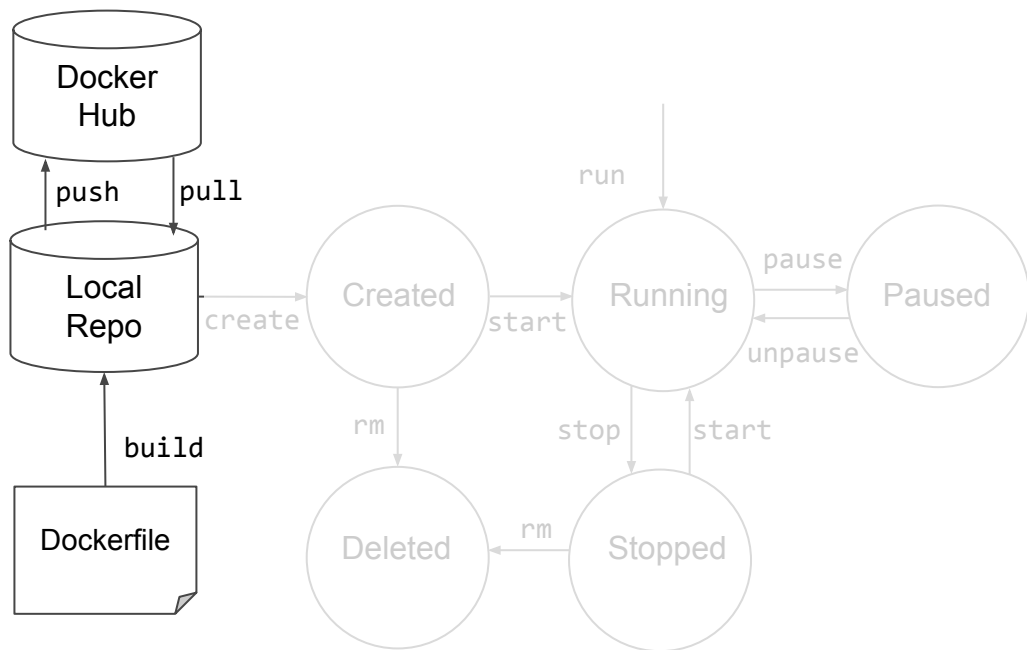
- ❑ **Immutable** template for containers
- ❑ Includes everything needed to run an application
 - ❑ code, runtime, system tools, system libraries, and settings

Container

- ❑ An instance of an image
- ❑ Add a **new writable layer** on top of the underlying image
 - ❑ all changes made to the running container (e.g., writing new files or modifying existing files) are written to this writable container layer

Docker container lifecycle

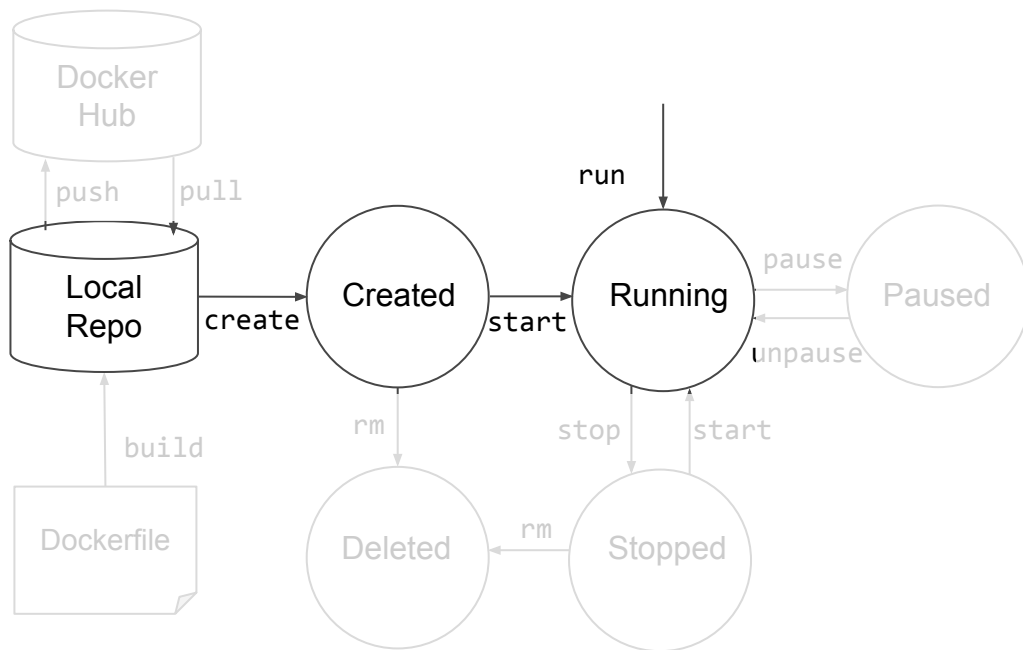
56



- ❑ Docker images can be pulled from a central repository (Docker Hub)
- ❑ Pulled images are saved in a local repository
- ❑ Custom images can be created and saved starting from a specific configuration file (Dockerfile)
- ❑ Custom images can be pulled to the central repository

Docker container lifecycle

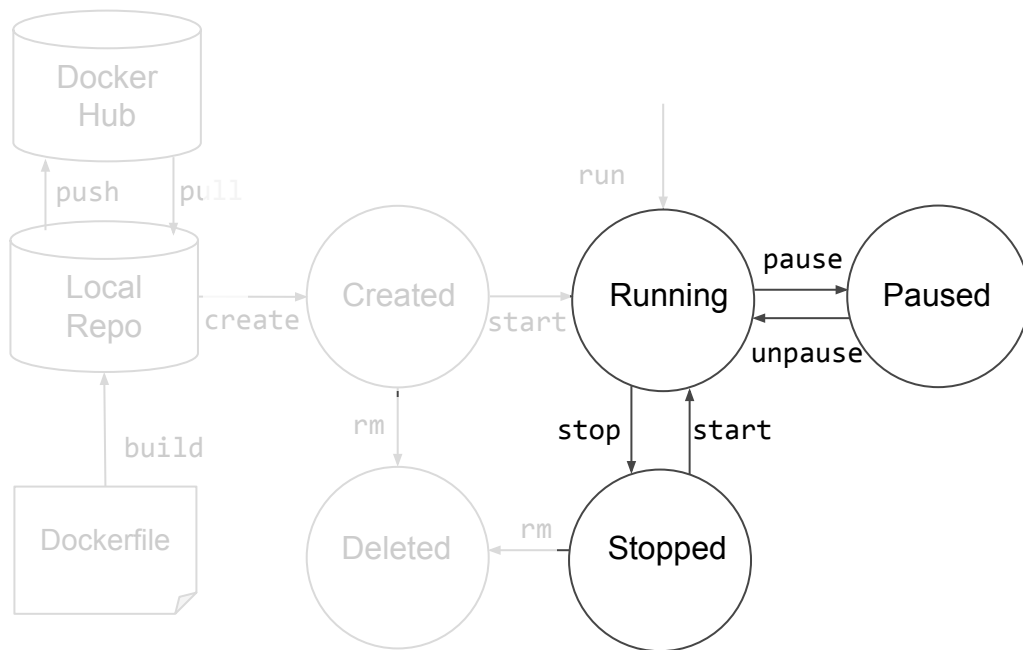
57



- A saved image can be used for creating a container (a writeable layer is added)
- Starting a container means running a default command contained in the image, namely the entrypoint (can be overridden)
- A container can be created and started using a single run command

Docker container lifecycle

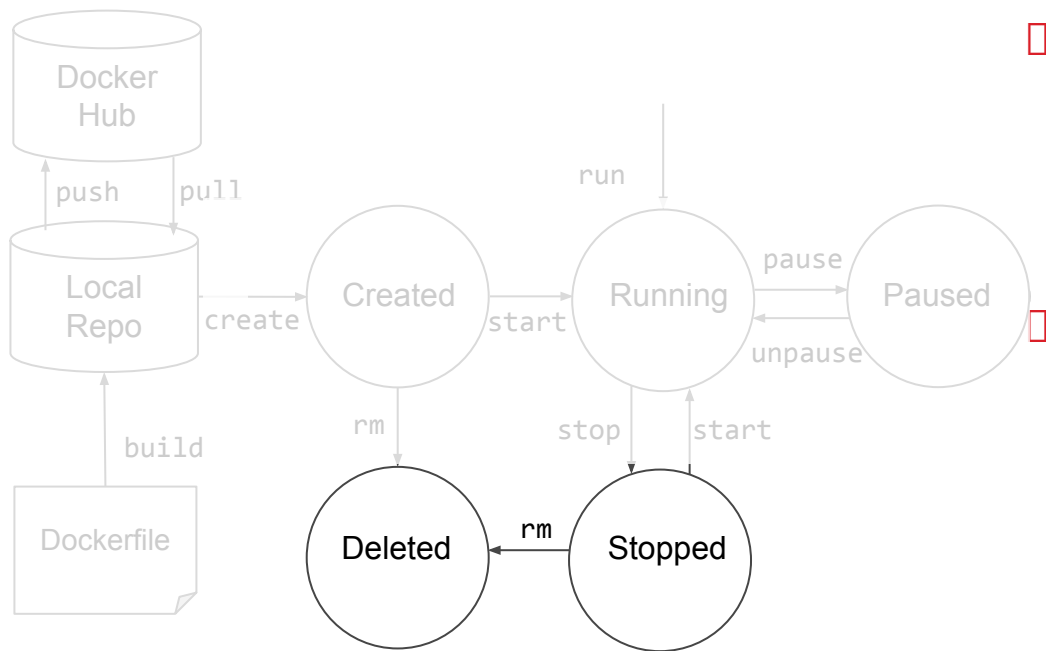
58



- After executing the entrypoint, the container stops
- A **foreground** process specified as the entrypoint can keep running the container
- A running container can be paused or stopped

Docker container lifecycle

59



- ❑ After a container is stopped, the writeable layer still exists
- ❑ Deleting a container permanently removes the associated writable layer

Docker lifecycle example: images

60

- ❑ Pull an image from the central hub
 - ❑ `docker pull <image>`
- ❑ Build an image from a Dockerfile
 - ❑ `docker build -t <image>`
- ❑ List images saved in the local repository
 - ❑ `docker images`
- ❑ Delete an image
 - ❑ `docker rmi <image>`

Docker lifecycle example: containers

61

- ❑ Start a container
 - ❑ `docker run <image>`
- ❑ Start a container overriding default entrypoint with `<newcmd>`
 - ❑ `docker run --entrypoint <newcmd> <image>`
- ❑ Execute a command `<cmd>` (e.g., `/bin/bash`) in a running container
 - ❑ `docker exec -it <containerID> <cmd>`
- ❑ Stop a container
 - ❑ `docker stop <containerID>`
- ❑ Remove a container
 - ❑ `docker rm <containerID>`
- ❑ List running and stopped containers
 - ❑ `docker ps -a`

Docker volumes

62

- ❑ Volumes can be used to save (persist) data and to share data between containers
- ❑ Volume is unrelated to the container layers: deleting a container does not involve deleting an associated volume
- ❑ A volume can be:
 - ❑ (anonymous/)named: managed internally by Docker itself
 - ❑ host: refers to a filesystem location of the host running Docker

Docker volumes: example

63

- ❑ Create a named volume
 - ❑ `docker volume create volumename`
- ❑ Run a container using the named volume
 - ❑ `docker run -v volumename:/path/in/container_filesystem`
- ❑ Running a container using a host volume
 - ❑ `docker run -v /path/on/host_filesystem:/path/in/container_filesystem`

Dockerfile

64

- ❑ Docker can build custom images automatically by reading the instructions from a Dockerfile
- ❑ Dockerfile is a text document that contains all the commands a user could call on the command line to assemble an image
- ❑ The *docker build* allows the execution of an automated build of an image starting from a Dockerfile

(Dockerfile reference: <https://docs.docker.com/engine/reference/builder/>)

Dockerfile example

65

Command	Description	Example
FROM <image>	Start building from this (base) image	FROM ubuntu
RUN <cmd>	Run the specified command	RUN apt install apache2
COPY <src> <dest>	Copy a file to the image fs	COPY vh.conf /etc/apache2/conf/
CMD ["exec", "param1", ...]	Configure the default command when container starts	CMD ["apache2", "-D", "FOREGROUND"]

Outline

66

- Docker images and containers
- **Docker networking**
- Docker compose

Docker networking

67

Docker supports different configurations, the two main ones being

▣ *Bridge* (default)

- ▣ isolated layer 3 networks enabling connected containers to communicate
- ▣ (can) allow the access to external networks masquerading connections with the host network configuration

▣ *Host*

- ▣ containers use the host network
- ▣ listening ports are exposed to the outside world

Docker networking: published ports

68

- A container connected to bridges is isolated and does not expose any of its ports to the outside world
- A published port can be made available to services running outside the container
- A published port is mapped to a port on the Docker host

Docker networking: examples

69

- ❑ Create network with a configured subnet and gateway address
 - ❑ `docker network create --driver bridge <networkname> --subnet=<ip/mask> --gateway=<ip/mask>`
- ❑ Connect a container to a network
 - ❑ `docker network connect <networkname> <containerid>`
- ❑ Run a container and expose a port
 - ❑ `docker run -h <host-name> -p <internal-port>:<exposed-port> --name <container-name> <image-name>`

Outline

70

- Docker images and containers
- Docker networking
- **Docker compose**

Docker compose

71

Compose is a tool for defining and running multi-container Docker applications

- A single file for providing configurations
- A single set of commands for configuring, building, and running all the containers

Docker compose configuration

72

- The Compose file (`docker-compose.yaml`) uses a standard, human-readable syntax, namely YAML* syntax. It defines
 - Services: configuration that is applied to each container (much like passing command-line parameters to *docker run*)
 - Networks (optional): define configuration of networks to be created
 - Volumes (optional): define configuration of volumes to be created

* <https://yaml.org/>

Docker compose configuration: example

73

services:

frontend container

app:

use a custom image

build:

directory containing Dockerfile

context: ./app

image name

image: custom_image

exposed ports (host:container)

ports:

- 8080:80

a mapped host volume

volumes:

- ./config/config.json:/etc/config.json

connected networks (defined in networks..)

networks:

ext:

ipv4_address: 192.168.100.100

int:

backend container

db:

pull an existing image

image: mariadb

a mapped named volume

volumes:

- db-content:/var/lib/mysql

networks:

int:

volumes:

db-content:

networks:

ext:

driver: bridge

ipam:

driver: default

config:

- subnet: 192.168.100.0/24

int:

driver: bridge

Docker compose: commands example

74

- ❑ (build images and) run services
 - ❑ `docker-compose up -d`
- ❑ stop services
 - ❑ `docker-compose stop`
- ❑ start services
 - ❑ `docker-compose start -d`
- ❑ stop and remove containers and networks
 - ❑ `docker-compose down`
- ❑ show logs (entrypoint output)
 - ❑ `docker-compose logs -f [service_name]`